

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import re

import plotnine
from plotnine import *
import matplotlib.pyplot as plt

import gc
import os.path
from os import path
import networkx

from research.data.sources import CachedDailyData, TickDB
from research.data.analysis import confusion

from research.functions import MultiFunction1D, SeriesFunction1D, TickFunction1D
from research.core.ggplot import ggcolors
from research.core.ggplot import scale_x_datetime_auto

pd.set_option('display.max_columns', 500)
pd.set_option('display.max_rows', 2000)
pd.options.display.max_colwidth = 100

plt.rcParams['figure.figsize'] = [8, 6]
plt.rcParams['figure.dpi'] = 100
```

1. older sites perform better and the impact outweighs other factors?
2. higher email sends lead to higher pvs. higher social leads to higher pvs?
3. certain words (clickbait or otherwise) lead to higher performance?
4. shorter stories outperform?
5. categories outperform, also that NY/CT vs NJ/PA have a bias on categories they report?
6. there is a sweet spot as to population that a given story should be sent to, however in general that most stories arent being sent to enough pop

Data Preparation

```
In [2]: def categoryFor(x):
    if "Police" in x:
        return "Police & Fire"
    if "Lifestyle" in x:
        return "Lifestyle"
    if "Schools" in x:
        return "Schools"
    if "Obituaries" in x:
        return "Obituaries"
    if "Business" in x:
        return "Business"
    if "Sports" in x:
        return "Sports"
    if "Weather" in x:
        return "Weather"
    if "Politics" in x:
        return "Politics"
```

```

if "Real Estate" in x:
    return "Real Estate"
else:
    return x

```

```

In [3]: df = pd.read_csv("../data/article_data_v3.csv", parse_dates=['publish_datetime'])
df["nviews_geo"] = df["total_pageviews"] / df["population_covered"]
df["primary_category"] = df["primary_category"].apply(categoryFor)
df["clickrate"] = (df["email_clicks"] / df["email_opens"]).fillna(0.0).replace([
df["openrate"] = (df["email_opens"] / df["email_sends"]).fillna(0.0).replace([np
df["email_geo"] = (df["email_sends"] / df["population_covered"]).fillna(0.0).rep

```

Basic Stats

% view by population size distribution

```

In [308... sns.distplot(df["nviews_geo"].apply(lambda x: min(x,0.25))).set_title('% views /


```

/usr/local/conda/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```

Out[308... Text(0.5, 1.0, '% views / population')

```



```
In [309... df["nviews_geo"].describe()
```

```
Out[309... count      25202.000000
mean         0.041171
std          0.162538
min          0.000000
25%          0.001906
50%          0.006984
75%          0.029476
max          10.525470
Name: nviews_geo, dtype: float64
```

Older Sites Perform Better than Newer?

```
In [275... age_relation = df.groupby("primary_site").agg({'primary_site_age': np.mean, 'nvi
```

```
In [276... pd.concat([
    df.query("primary_site_age < 365").nviews_geo.describe(),
    df.query("primary_site_age >= 365 and primary_site_age < 1500").nviews_geo.d
    df.query("primary_site_age >= 1500").nviews_geo.describe(),
], keys=['< 1 yr', '1-2 yrs', '> 2 yrs'], axis=1)
```

```
Out[276...      < 1 yr      1-2 yrs      > 2 yrs
count  7469.000000  7202.000000  10531.000000
mean     0.006472    0.007365    0.088901
std     0.035400    0.015174    0.241385
min     0.000000    0.000000    0.000000
25%     0.000739    0.001381    0.012972
50%     0.002162    0.003196    0.034794
75%     0.005502    0.007536    0.088224
max     2.594834    0.475761    10.525470
```


Higher emails lead to more story page views?

- there is a postive correlation between emails and increased story views (though only 43%)

```
In [277... sns.jointplot(df["email_sends"], df["email"])
```

```
/usr/local/conda/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWa
rning: Pass the following variables as keyword args: x, y. From version 0.12, th
e only valid positional argument will be `data`, and passing other arguments wit
hout an explicit keyword will result in an error or misinterpretation.
```

```
Out[277... <seaborn.axisgrid.JointGrid at 0x7fd30e9c18b0>
```



```
In [278...] df[["email_sends","email"]].corr()
```

```
Out[278...]
      email_sends  email
email_sends    1.000000  0.428275
email          0.428275  1.000000
```

```
In [279...] df[["email_sends","email_opens"]].corr()
```

```
Out[279...]
      email_sends  email_opens
email_sends    1.000000    0.59347
email_opens    0.59347    1.00000
```

```
In [280...] df[["email_sends","total_pageviews"]].corr()
```

```
Out[280...]
      email_sends  total_pageviews
email_sends    1.000000    0.066597
```

	email_sends	total_pageviews
total_pageviews	0.066597	1.000000


Do shorter stories outperform?

- body length **does not** show any relationship to page view performance of story (3% correlation)
- this makes sense, as people will view a story based on title and then decide how much to read, as the story progresses

```
In [324...] sns.distplot(df["body_length"].apply(lambda x: min(x,20000))).set_title("distrib
```

```
/usr/local/conda/lib/python3.8/site-packages/seaborn/distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
Out[324...] Text(0.5, 1.0, 'distribution of body length')
```




```
In [326...] sns.jointplot(df["body_length"],df["nviews_geo"].apply(lambda x: min(x,1)))
```

```
/usr/local/conda/lib/python3.8/site-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
<seaborn.axisgrid.JointGrid at 0x7fd30fb560d0>
```

Out[326...



In [283... `df[["body_length","nviews_geo"]].corr()`

Out[283...

	body_length	nviews_geo
body_length	1.000000	0.035623
nviews_geo	0.035623	1.000000

Which categories outperform, by geo?

- top absolute # views involve police and fire (but those are the most frequent stories)
- in terms of % views / population, police and fire range below the top 5 in most states

In [284... `agg = df.groupby(["primary_category"]).nviews_geo.agg(["mean", "count"])`
`agg.query("count > 5").sort_values("mean", ascending=False)`

Out[284... `mean count`

primary_category	mean	count
primary_category		
Schools	0.066937	757
News	0.058629	6517
Events	0.054936	10
Business	0.052768	1123
Obituaries	0.047849	478
Lifestyle	0.040955	1420
Police & Fire	0.031830	13523
Weather	0.031269	451
Real Estate	0.025130	209
Sports	0.021126	135
Politics	0.020616	413
Neighbors	0.016074	8
Traffic	0.014027	82
News, Politics	0.010168	53
News, Traffic	0.003853	13

In [285...

```
agg = df.groupby(["primary_state", "primary_category"]).nviews_geo.agg(["mean",
agg.sort_values(["primary_state", "mean"], ascending=False)
```

Out[285...

		mean	count
primary_state primary_category			
PA	Weather	0.010223	20
	Obituaries	0.005713	73
	Police & Fire	0.002273	1230
	News	0.002255	141
	Lifestyle	0.001876	75
	Schools	0.001742	22
	Business	0.001630	47
	Real Estate	0.001171	6
	News, Politics	0.000664	12
NY	Business	0.094524	259
	Obituaries	0.091040	144
	Schools	0.086713	301
	News	0.076249	3597

		mean	count
primary_state	primary_category		
NJ	Lifestyle	0.046845	648
	Weather	0.032048	223
	Police & Fire	0.024001	3682
	Traffic	0.020263	46
	Politics	0.020163	325
	Sports	0.018006	77
	Real Estate	0.016905	162
	Schools	0.081630	257
	Politics	0.074791	15
	Business	0.070190	432
	Real Estate	0.068099	29
	Events	0.059771	9
	News	0.058133	1272
	Lifestyle	0.055405	380
	Police & Fire	0.046375	6607
	Obituaries	0.037823	230
	Sports	0.034890	37
	Weather	0.024937	49
	News, Politics	0.014973	35
	Traffic	0.005814	13
MA	News, Traffic	0.004618	10
	News	0.002453	369
	Business	0.002356	166
	Lifestyle	0.002204	87
	Police & Fire	0.001926	510
	Schools	0.001261	25
	News, Politics	0.001142	6
	Weather	0.000953	32
	Traffic	0.000435	9
CT	Real Estate	0.052768	10
	Weather	0.043297	127
	News	0.028693	1138
	Lifestyle	0.027884	230
	Schools	0.023168	152

	mean	count
primary_state primary_category		
Obituaries	0.023010	27
Police & Fire	0.021340	1494
Business	0.018204	219
Traffic	0.015193	9
Politics	0.012226	68
Sports	0.011072	15

Headlines that drive views

```
In [315... dropwords = [
    'the', 'for', 'at', 'as', 'of', 'with', 'are', 'was', 'were', 'is', 'a', 'an',
    'you', 'him', 'her', 'in', 'to', 'be', 'from', 'on', 'heres', 'say', 'two', "your",
    'connecticut', 'patterson', 'paramus', 'jersey', 'westchester', 'hudson', 'valley',
    'newark', 'rockland', 'long', 'island', 'year', 'route', "overpeck", "ridegwood",
    'northern', 'nj', 'ny', 'ct', 'morris',
    '0', '1', '2', '3', '4', '5', '6', '7', '8', '9',
    '10', '11', '12', '13', '14', '15', '16', '19',
    '20', '21', '22', '23', '24', '25', '26', '27', '28', '29',
    '30', '31', '32', '33', '34', '35', '36', '37', '38', '39',
    '40', '41', '42', '43', '44', '45', '46', '47', '48', '49',
    '50', '51', '52', '53', '54', '55', '56', '57', '58', '59',
]

def wordsFor (row):
    stamp = row["publish_datetime"]
    title = str(row["title"])
    state = str(row["primary_state"])
    category = str(row["primary_category"])
    score = float(row["nviews_geo"])
    words = [x.lower() for x in re.findall(r'\w+', title.replace("'", ""))]
    filtered = [x for x in words if not x in dropwords]
    return pd.Series([stamp, state, category, filtered, score], index=['stamp', 'st

wordv = df.apply(wordsFor, axis=1)
```

```
In [287... def filtered (query = "score > 0", by = 'freq'):
    freqs = {}
    counts = {}
    sub = wordv.query (query)
    for i in range(sub.shape[0]):
        words = sub.iloc[i].words
        score = float(sub.iloc[i].score)
        if by == "freq" or np.isnan(score):
            weight = 1
        else:
            weight = int(score * 100)

        for word in words:
            if word in freqs:
                freqs[word] += weight
```

```

        counts[word] += 1
    else:
        freqs[word] = weight
        counts[word] = 1

    if by == "mean1":
        for word in counts.keys():
            freqs[word] = freqs[word] / counts[word]
    elif by == "mean2":
        for word in counts.keys():
            freqs[word] = freqs[word] / np.sqrt(counts[word])

    return freqs

```

Overall word cloud

In [288...

```

allwords = []
for i in range(wordv.shape[0]):
    words = wordv.iloc[i].words
    allwords += words

```

In [289...


```

from wordcloud import WordCloud

# Create the wordcloud object
wordcloud = WordCloud(width=480, height=480, margin=0).generate(" ".join(allword

# Display the generated image:
fig = plt.figure(figsize=(12,8), dpi= 100)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.margins(x=0, y=0)
plt.show()

```




Performance Adjusted: By Cum % of Views

```
In [290... freqs = filtered("score > 0", by='mean1')

# Create the wordcloud object
wordcloud = WordCloud(width=480, height=480, margin=0).generate_from_frequencies

# Display the generated image:
fig = plt.figure(figsize=(12,8), dpi= 100)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.margins(x=0, y=0)
plt.show()
```

Differences between NY/CT and NJ/PA

```
In [320... sub = df.query("primary_state != 'MA'").copy()
sub["area"] = sub.primary_state.apply(lambda x: "PANJ" if x == "PA" or x == "NJ
```

Click rates

```
In [321... sub.groupby("primary_state").clickrate.mean()
```

```
Out[321... primary_state
CT      0.167658
NJ      0.161063
NY      0.176344
PA      0.125623
Name: clickrate, dtype: float64
```

```
In [322... sub.groupby("primary state").openrate.mean()
```


		mean	count
area	primary_category		
	News	0.064819	4735
	Business	0.059557	478
	Lifestyle	0.041878	878
	Weather	0.036130	350
	Police & Fire	0.023233	5176
	Traffic	0.019434	55
	Real Estate	0.018990	172
	Politics	0.018790	393
	Sports	0.016876	92

In [310...

```
agg = sub.groupby(["area", "primary_category"]).total_pageviews.agg(["sum", "count"])
agg.sort_values(["area", "sum"], ascending=False)
```

Out[310...

		sum	count
area	primary_category		
PANJ	Police & Fire	43286162	7838
	News	8660473	1413
	Lifestyle	3444741	455
	Business	3314769	479
	Schools	2742639	279
	Obituaries	1882894	303
	Weather	299016	69
	Real Estate	235159	35
	Sports	161614	38
	News, Politics	113896	47
	Politics	51789	16
	News, Traffic	43186	11
	Events	32225	9
	Traffic	30142	18
NYCT	News	43974779	4735
	Police & Fire	16635824	5176
	Weather	11327161	350
	Lifestyle	4567082	883
	Business	3019377	478
	Schools	3008611	473

		sum	count
area	primary_category		
	Politics	1644665	406
	Obituaries	1005842	171
	Real Estate	350702	172
	Sports	235643	92
	Traffic	207747	55

In []:

In []: